

Introduction to JAGS and rJAGS

Athena Chen

Tuesday, April 28, 2020

JHU Biostatistics Student Computing Club

Just Another Gibbs Sampler (JAGS)

- Introduction, installation, and useful resources
- Components of a JAGS program
- The JAGS language
- rJAGS

Example: Normal-Normal Model

Example: Mixture of Normal Distributions

- The label switching problem

Additional Tricks and Tips

- Custom likelihoods: the zero's and one's trick

Just Another Gibbs Sampler (JAGS)

Just Another Gibbs Sampler (JAGS) is a program for conducting Bayesian analyses and modeling using MCMC.

- Written in C++ and available on all major OS.
- Slice sampler (Neal 2003, *The Annals of Statistics*)

Installation

- **Sourceforge**: <https://sourceforge.net/projects/mcmc-jags>
- **Homebrew** (MacOS): `brew install jags`

Useful resources

- [User manual](#)

The JAGS language is very friendly to R users.

Relations

deterministic assignment	<code>mu[i] <- alpha + beta*x[i]</code>
stochastic assignment	<code>Y[i] ~ dnorm(mu[i], tau)</code>

Arrays and subsetting¹

vectors	<code>mu[1], mu[1:10]</code>
arrays	<code>B[r,c], B[1:M,c], B[1:M,1:N]</code>

Vector-construction and for-loops

Vector construction	<code>y <- c(x1, x2)</code>
for-loops	<code>for(i in 1:N) { ... }</code>

¹Indices cannot be stochastic or repeated.

Components of a JAGS Program

1. Model definition
2. Compilation of model into computer memory
3. Initialization of model
4. Adaptation and burn-in
5. Monitoring or saving MCMC-generated values

Example: Normal-Normal Model

Normal-Normal model

Let $N(\mu, \tau^2)$ denote a normal distribution with mean μ and precision τ^2 .

Suppose

$$Y_i | \mu, \tau^2 \stackrel{iid}{\sim} N(\mu, \tau^2) \quad \text{likelihood}$$

$$\mu \sim N(\mu_0, \tau_0^2) \quad \text{prior}$$

$$\tau^2 \sim \text{Gamma}(a, b)$$

Then we know that,

$$\mu | \tau^2, \mathbf{Y} \sim N \left(\frac{\tau_0^2 \mu_0 + n \tau^2 \bar{Y}}{\tau_0^2 + n \tau^2}, \tau_0^2 + n \tau^2 \right)$$

$$\tau^2 | \mu, \mathbf{Y} \sim \text{Gamma} \left(a + \frac{n}{2}, b + \frac{1}{2} \sum_i (Y_i - \mu)^2 \right)$$


Simulate data

```
mu <- 0; tau2 <- 0.01           # Define simulation parameters

set.seed(2020428)

n <- 100                        # Number of observations in our data
Y <- rnorm(n, mu, sqrt(1/tau2)) # Simulate data
```

Distribution of simulated data



We define the model in the file `nn_model.jags`.

```
model {  
  mu ~ dnorm(mu_0, tau2_0)  
  tau2 ~ dgamma(a, b)  
  
  for(i in 1:n){  
    Y[i] ~ dnorm(mu, tau2)  
  }  
}
```

Define priors

```
library(rjags)                # R interface to JAGS
library(ggmcmc)               # tidy output objects

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg      ggplot2

# Define prior parameters
mu_0 <- 0; tau2_0 <- 0.01     # Prior parameters for mu
a <- 1; b <- 0.1              # Prior parameters for tau

# Define data
data_list <- list(mu_0 = mu_0,
                  tau2_0 = tau2_0,
                  a = a, b = b,
                  n = n, Y = Y)
```

Compile model into memory and run model

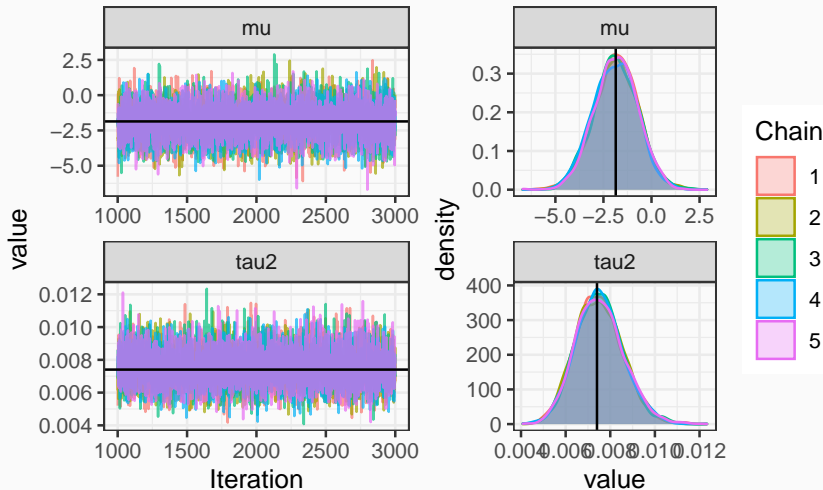
```
# Compile model into memory
jags_model <- jags.model("./model/nn_model.bugs",
                        n.chains = 5,
                        data = data_list,
                        quiet = TRUE)

# burn-in of 1000 samples
update(jags_model, 1000)


# draw 2000 and convert samples to tidy format
samples <- coda.samples(jags_model,
                        variable.names = c("mu", "tau2"),
                        n.iter = 2000) %>%

ggs()
```

Posterior trace and density plots



Chain autocorrelation



Example: Finite Mixture of Normal Distributions

Suppose we have data that comes from a mixture of normal distributions:

$$Z_i | \pi \sim \text{Bernoulli}(\pi)$$

$$Y_i | Z_i, \mu_1, \mu_2, \tau_2^2 \stackrel{iid}{\sim} \begin{cases} N(\mu_1, \tau^2) & \text{if } Z_i = 1 \\ N(\mu_2, \tau^2) & \text{if } Z_i = 0 \end{cases}$$

Let,

$$\mu_i \sim N(\mu_{0i}, \tau_0^2)$$

$$\tau^2 \sim \text{Gamma}(a, b)$$

$$\pi \sim \text{Beta}(\alpha, \beta)$$

Simulate data


```
mu_1 <- -2.75; mu_2 <- 2.75; tau2 <- 1; pi <- 0.4
```

```
set.seed(689934)
```

```
n <- 1000
```

```
Z <- rbinom(n, 1, pi)
```

```
Y <- Z*rnorm(1000, mu_1, sqrt(1/tau2)) +  
  (1-Z)*rnorm(1000, mu_2, sqrt(1/tau2))
```



Define model

We define the model in the file `mixed_nn_model.jags`.

```
model {  
  mu_1 ~ dnorm(mu_01, tau2_0)  
  mu_2 ~ dnorm(mu_02, tau2_0)  
  tau2 ~ dgamma(a, b)  
  pi ~ dbeta(alpha, beta)  
  
  for(i in 1:n){  
    Z[i] ~ dbern(pi)  
    mu[i] <- Z[i]*mu_1 + (1-Z[i])*mu_2  
    Y[i] ~ dnorm(mu[i], tau2)  
  }  
}
```

Note that in JAGS, we can explicitly model the class label, Z_i .

Define priors

```
# Define priors for mu
mu_01 <- -10; mu_02 <- 10;
tau2_0 <- 0.01

# Prior parameters for tau
a <- 1; b <- 0.01

# Prior parameters for pi
alpha <- 4; beta <- 6

# Define data
data_list <- list(mu_01 = mu_01, mu_02 = mu_02,
                  tau2_0 = tau2_0,
                  a = a, b = b,
                  alpha = alpha, beta = beta,
                  n = n, Y = Y)
```

Compile model into memory and run model


```
# Compile model into memory
jags_model <- jags.model("./model/mixed_nn_model.bugs",
                        n.chains = 5,
                        data = data_list,
                        quiet = TRUE)

# burn-in of 1000 samples
update(jags_model, 1000)


# draw 2000 and convert samples in tidy format
samples <- coda.samples(jags_model,
                        variable.names = c("mu_1", "mu_2", "pi",
                                           "tau2", "Z"),
                        n.iter = 2000) %>%

ggs()
```


Posterior trace and density plots



Posterior probabilities of group 1 vs group 2



Chain autocorrelation



Non-identifiability and the label-switching problem

One of the main problems with mixture models is an identifiability problem often referred to as the **label-switching problem**.

Briefly, the label switching problem arises when the prior and likelihood are invariant under permutation of class labels.

Next week, **Tuesday, May 5th**, Jacob will implement this mixture model in **Stan** and discuss the label-switching problem and how to address it in more depth.

Additional Tips and Tricks

Custom likelihoods: the Zero's Trick

The **zero's trick** is an approach to specifying custom likelihoods in JAGS.

Suppose L is the custom likelihood we would like to use. The probability of only observing zeros for a $\text{Poisson}(\lambda)$ distribution is given by,

$$f(0|\lambda) = \frac{\lambda^0 e^{-\lambda}}{0!} = e^{-\lambda}$$

Thus, if we define $\lambda = -\log(L)$, $f(0|\lambda) = e^{-(-\log(L))} = L$, and we can obtain the correct likelihood contribution.

```
model{  
  # rest of model block  
  for(i in 1:N){  
    zeros[i] <- 0                # Observation of all zeros  
    zeros[i] ~ dpois(lambda[i])  # Specify poisson dist  
    lambda[i] <- -log(L)         # Define custom likelihood  
  }  
}
```

Custom likelihoods: the One's Trick

An alternative to the zero's trick is the **one's trick**. Again, if L is the custom likelihood we would like to use, the probability of only observing ones for a $\text{Bernoulli}(\pi)$ distribution is given by,

$$f(1|\pi) = \pi^1(1 - \pi)^0 = \pi$$

Thus, if we define $\pi = L$, $f(0|\pi) = L$, and we can obtain the correct likelihood contribution.

```
model{  
  # rest of model block  
  for(i in 1:N){  
    ones[i] <- 1           # Observation of all ones  
    ones[i] ~ dbern(ones[i]) # Specify bernoulli dist  
    pi[i] <- L             # Define custom likelihood  
  }  
}
```